

Thrust OpenMP Backend


OpenMP Backend on Intel Westmere Xeon Processors

Compile (Gaussian Fit)

- Compiling with CUDA Backend
 - `nvcc -o a.out -arch=sm_20 solution2.cu -L./rootstuff/ -lRootUtils`
- Compiling with OpenMP Backend (icc and g++)
 - `g++ -O2 -o omp.out solution2.cpp -fopenmp -DTHRUST_DEVICE_BACKEND=THRUST_DEVICE_BACKEND_OMP -lgomp -I $CUDA_HOME/include -L./rootstuff/ -lRootUtils`
 - `icc -O2 -o omp.out solution2.cpp -fopenmp -DTHRUST_DEVICE_BACKEND=THRUST_DEVICE_BACKEND_OMP -lgomp -I $CUDA_HOME/include -L./rootstuff/ -lRootUtils`


Gauss Fit: CUDA vs OpenMP (12 Threads)

Gauss Fit Backends (12 Threads)




Gauss Fit: CUDA vs OpenMP (6 Threads)

Gauss Fit Backends (6 Threads)




Gauss Fit: CUDA vs OpenMP (2 Threads)

Gauss Fit Backends (2 Threads)

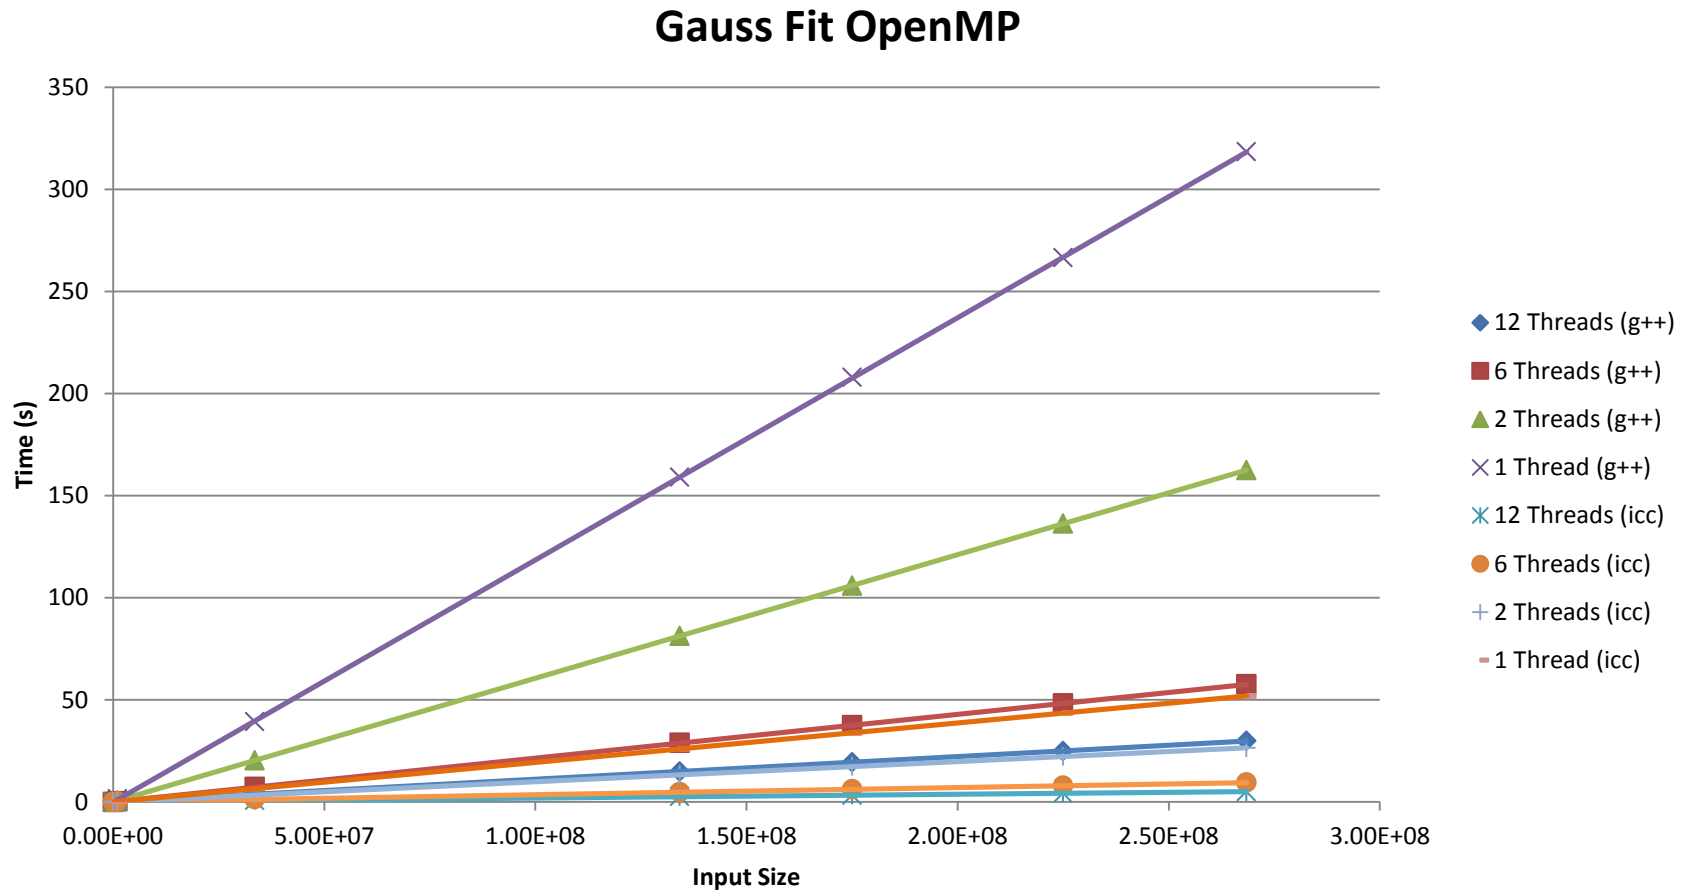


Gauss Fit: CUDA vs OpenMP (1 Thread)

Gauss Fit Backends (1 Thread)



Gauss Fit: Number of OpenMP Threads




Compile (Chi-square Fit)

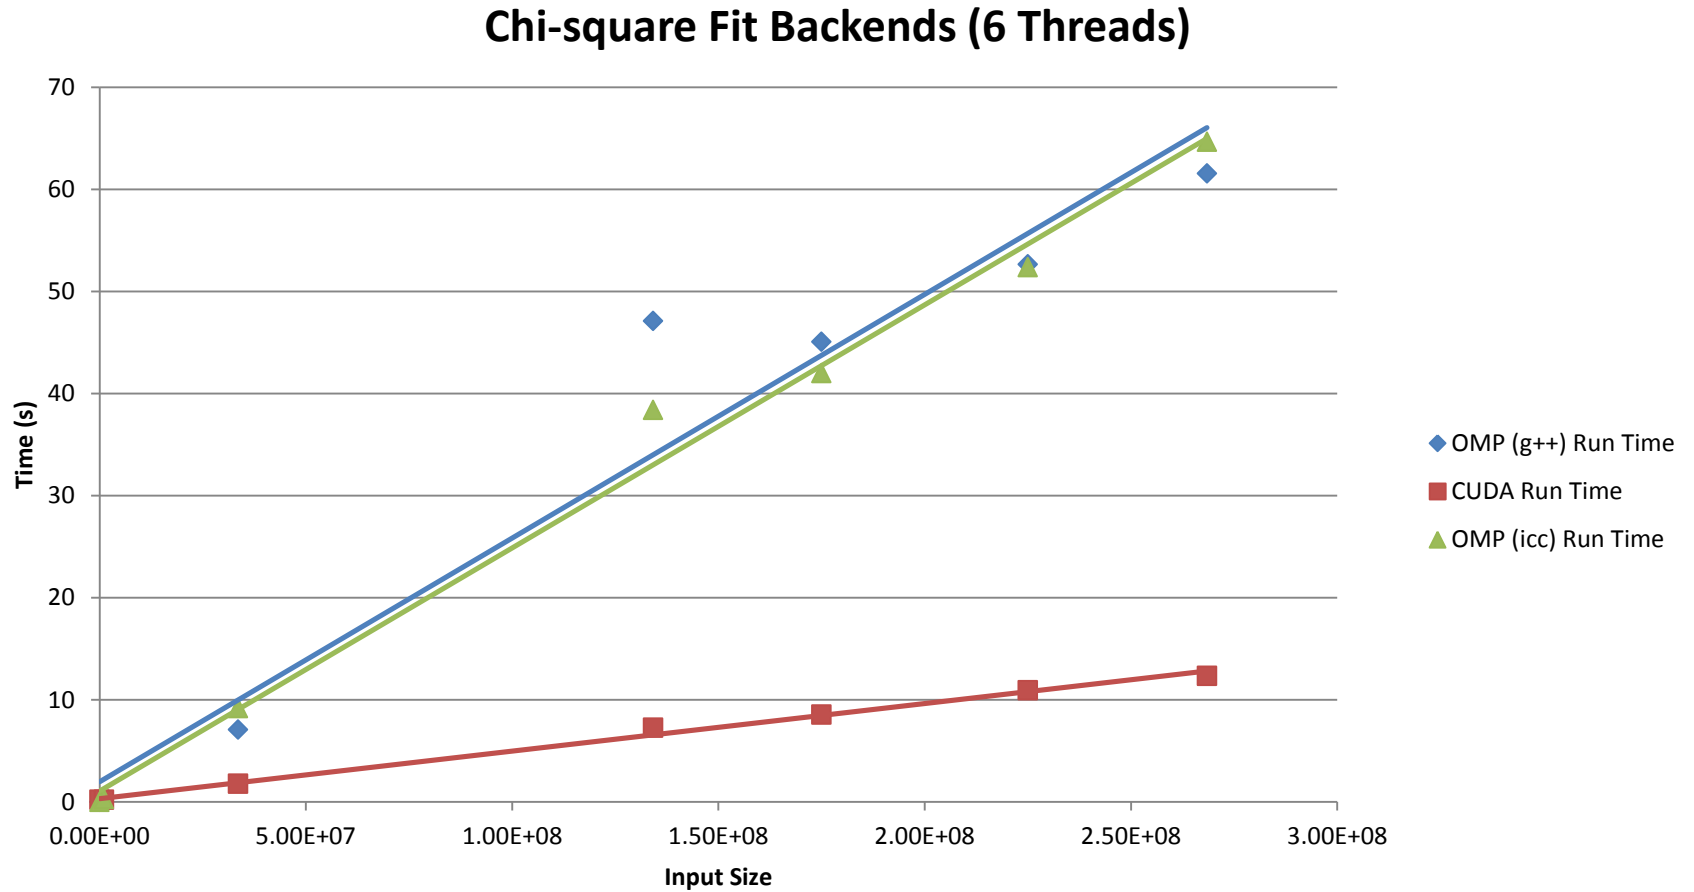
- Compiling with CUDA Backend
 - `nvcc -o a.out -arch=sm_20 example2.cu -L./rootstuff/ -lRootUtils`
- Compiling with OpenMP Backend
 - `g++ -O2 -o omp.out example2.cpp -fopenmp -DTHRUST_DEVICE_BACKEND=THRUST_DEVICE_BACKEND_OMP -lgomp -I $CUDA_HOME/include -L./rootstuff/ -lRootUtils`
 - `icc example2.cpp -o omp.out -fopenmp -DTHRUST_DEVICE_BACKEND=THRUST_DEVICE_BACKEND_OMP -lgomp -I $CUDA_HOME/include -O2 -L./rootstuff/ -lRootUtils`

Chi-square Fit: CUDA vs OpenMP (12 Threads)

Chi-square Fit Backends (12 Threads)




Chi-square Fit: CUDA vs OpenMP (6 Threads)




Chi-square Fit: CUDA vs OpenMP (2 Threads)

Chi-square Fit Backends (2 Threads)



Chi-square Fit: CUDA vs OpenMP (1 Thread)


Chi-square Fit Backends (1 Thread)



Chi-square Fit: Number of OpenMP Threads

(icc not included because times were similar)

Chi-square Fit OpenMP




OpenMP Backend on Intel Sandybridge Xeon processors

(Chi-square fit algorithm optimized)


Gauss Fit: CUDA vs. OpenMP (16 Threads)

Gauss fit Backends (16 Threads)




Gauss Fit: CUDA vs. OpenMP (12 Threads)

Gauss fit Backends (12 Threads)




Gauss Fit: CUDA vs. OpenMP (6 Threads)

Gauss fit Backends (6 Threads)




Gauss Fit: CUDA vs. OpenMP (2 Threads)

Gauss fit Backends (2 Threads)




Gauss Fit: CUDA vs. OpenMP (1 Threads)

Gauss fit Backends (1 Thread)

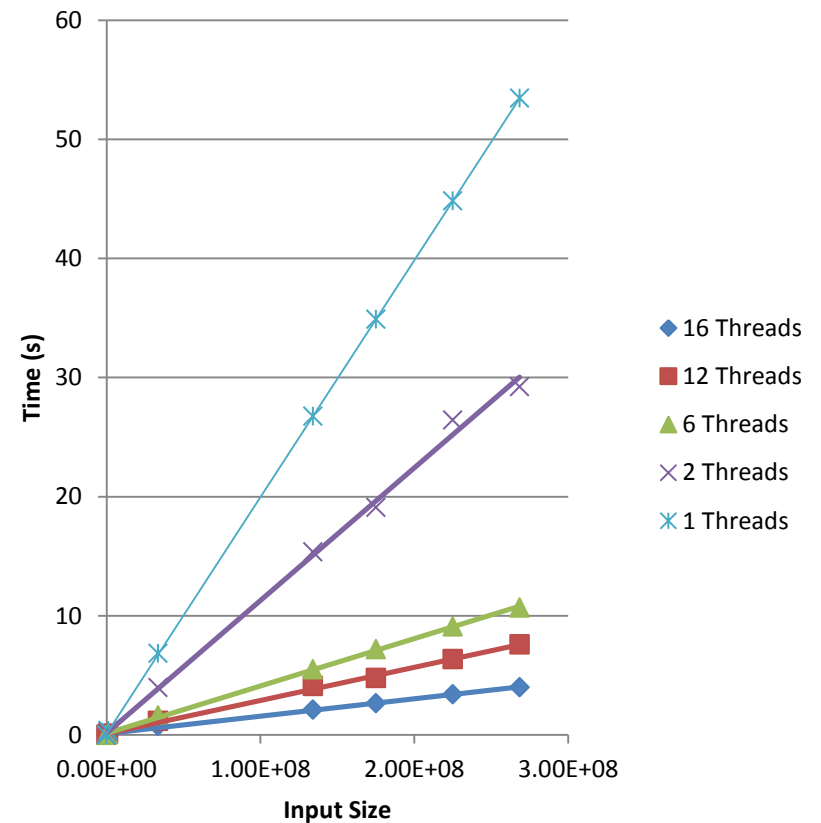


Gauss Fit: Number of OpenMP Threads

OpenMP Chi-square fit (g++)




OpenMP Chi-square fit (icc)



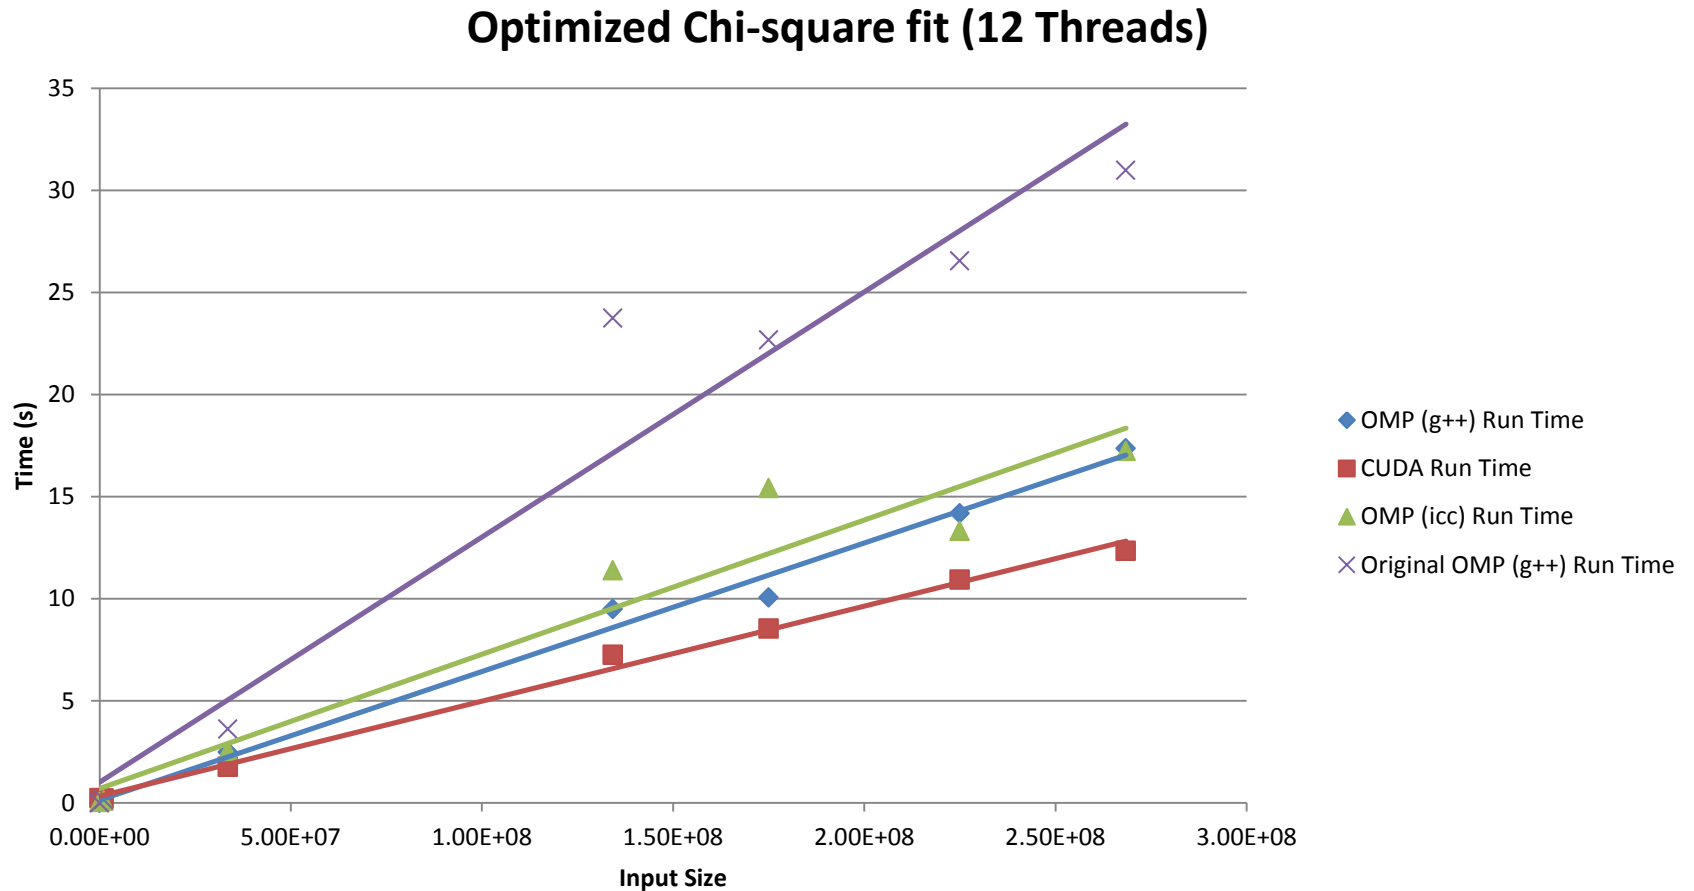
Optimizing Chi-square fit algorithm

- Instead of executing $\text{chisq} /= \text{error}$ in each loop, $\text{chisq} *= \text{oneOverError}$ is executed in each loop
- oneOverError is set equal to $1/\text{error}$ as an argument to the functor and therefore is only calculated once.


Optimized Chi-square fit: CUDA vs. OpenMP (16 Threads)



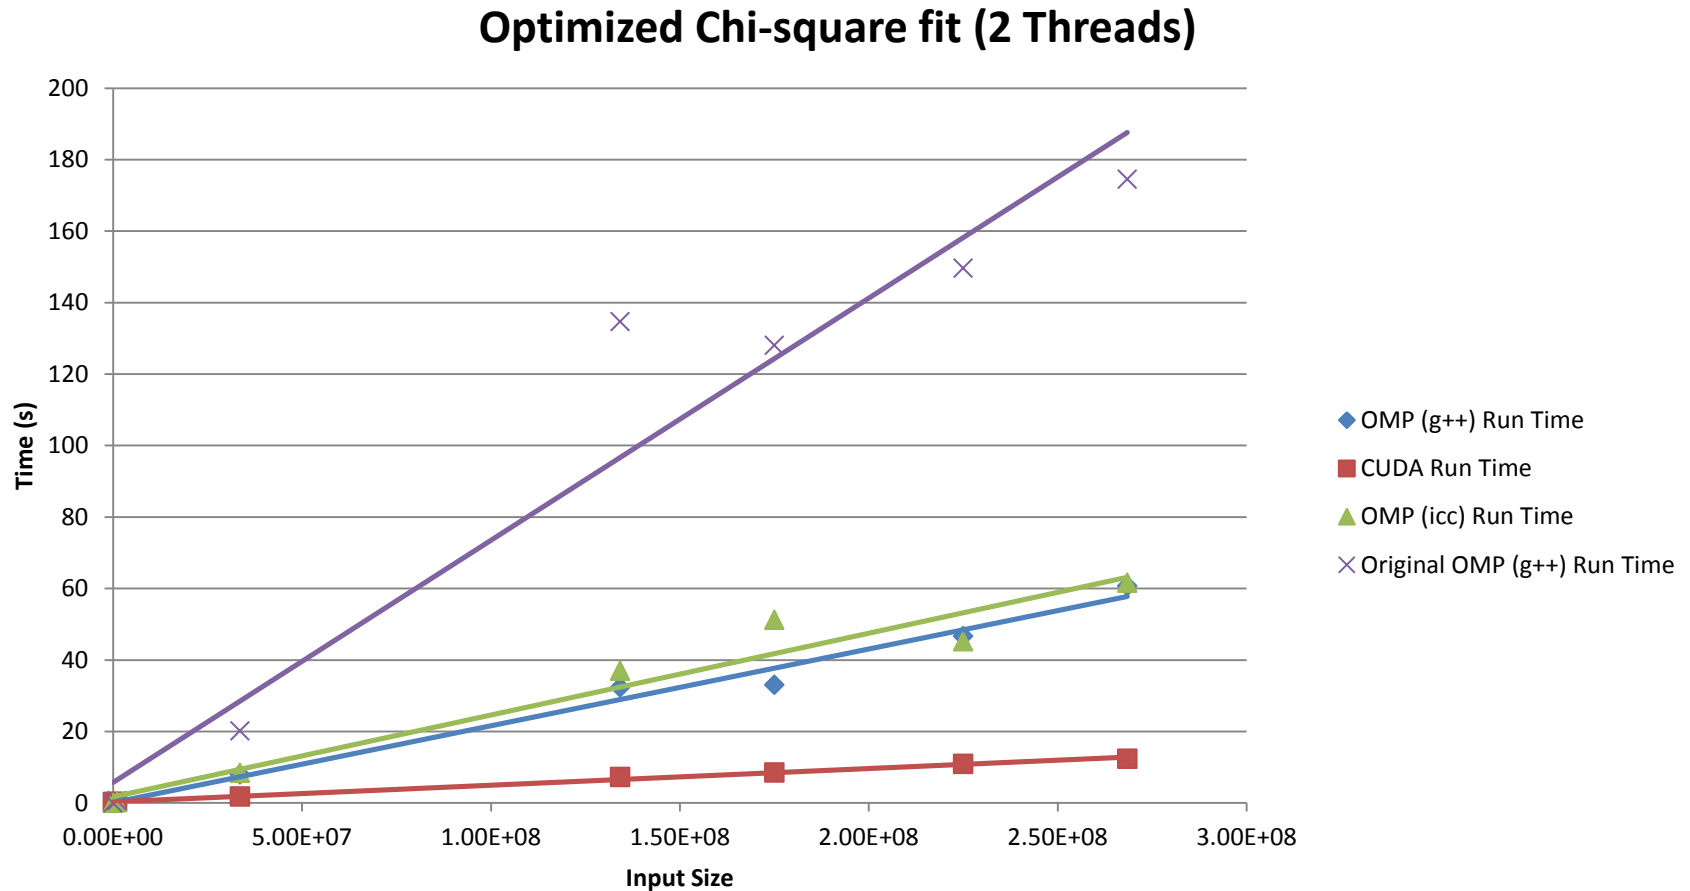
Optimized Chi-square fit: CUDA vs. OpenMP (12 Threads)



Optimized Chi-square fit: CUDA vs. OpenMP (6 Threads)




Optimized Chi-square fit: CUDA vs. OpenMP (2 Threads)




Optimized Chi-square fit: CUDA vs. OpenMP (1 Thread)

Optimized Chi-square fit (1 Thread)



Optimized Chi-square Fit: Number of OpenMP Threads



GNU Compiler vs. Intel Compiler

GNU Compiler vs. Intel Compiler

- The Gaussian fit example runs faster when compiled with the Intel compiler because the Intel compiler efficiently vectorizes the code.
- No performance difference is seen between the Chi-square fit example with the Intel compiler and the GNU compiler because the Chi-square fit example is memory bound.
- A hybrid code with the data layout of the Chi-square fit example and the mathematical work of the Gaussian fit example was used to test that the Chi-square fit example is memory bound.
- Due to the extra mathematical work, the hybrid code runs faster when compiled with the Intel compiler due to vectorization.

GNU Compiler vs. Intel Compiler (cont.)

```
__device__ double operator () (tuple<double, double> xypair)
{
    // Extract x and y from tuple
    double xval = get<0>(xypair);
    double yval = get<1>(xypair);

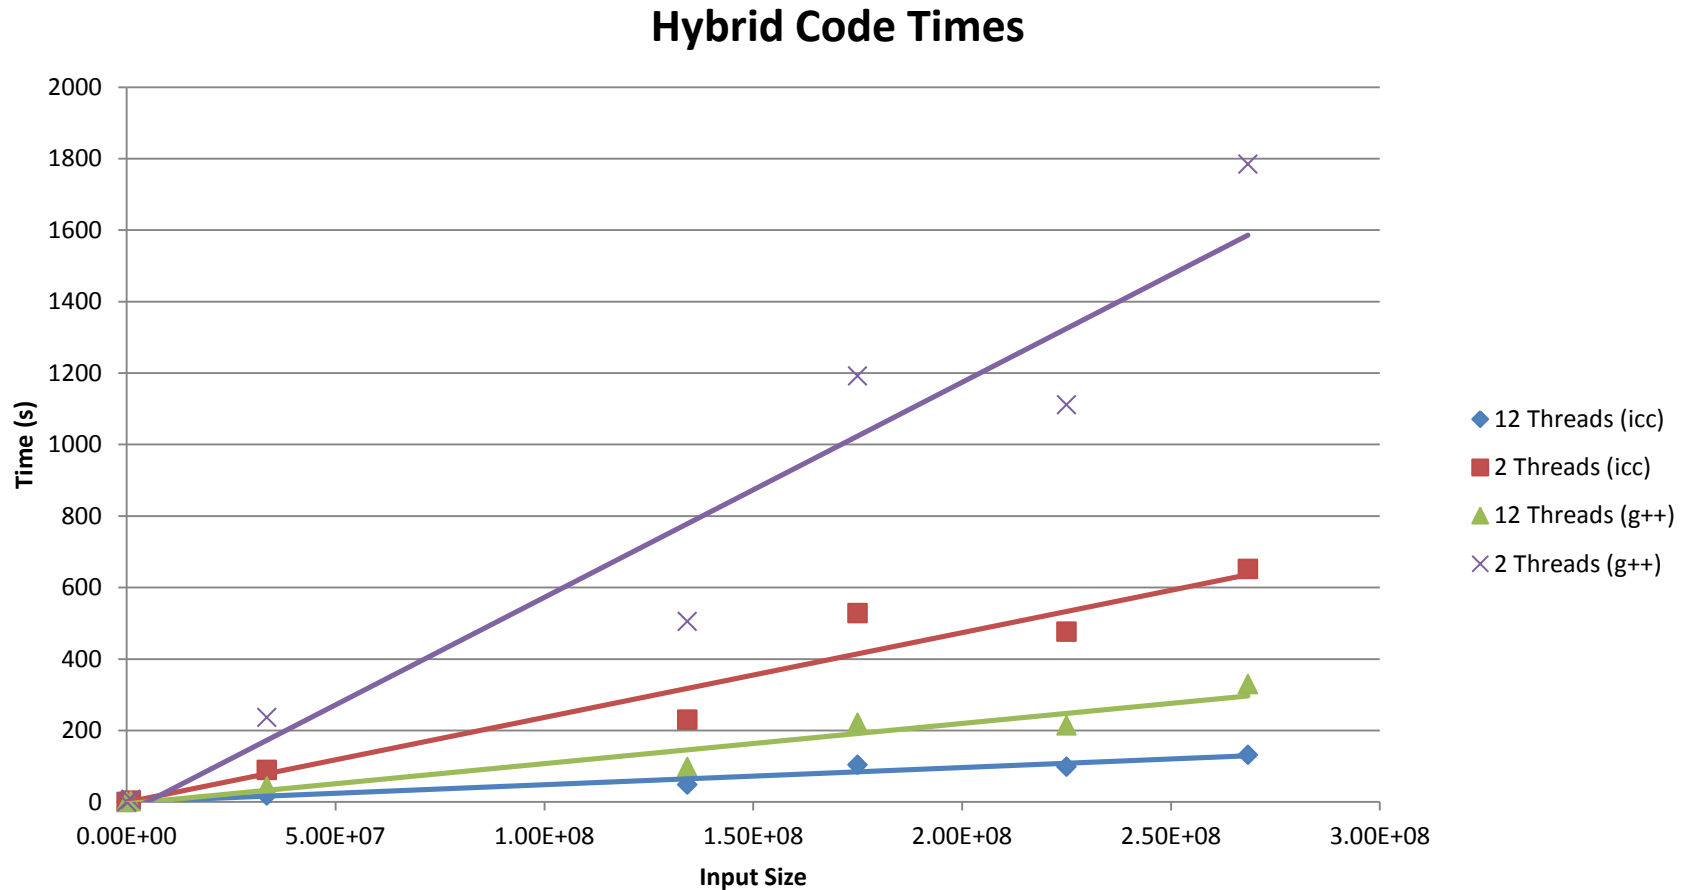
    // Calculate Gauss for x
    double retX = (xval - mean);
    retX /= sigma;
    retX *= retX;
    retX = exp(-0.5*retX);
    retX /= sigma;
    retX /= sqrt(2*M_PI);

    // Calculate Gauss for y
    double retY = (yval - mean);
    retY /= sigma;
    retY *= retY;
    retY = exp(-0.5*retY);
    retY /= sigma;
    retY /= sqrt(2*M_PI);
    return (0.5 * retX) + (0.5 * retY); // Return average
}
```

```
void chisq (int& npar, double* deriv, double& fVal,
           double param[], int flag) {
    double mean = param[0];
    double sigma = param[1];
    ChisqFunctor functor(mean, sigma);
    // Note hardcoded error!

    double initVal = 0;
    fVal = transform_reduce(
        make_zip_iterator(make_tuple(
            dev_xvals->begin(),
            dev_yvals->begin())),
        make_zip_iterator(make_tuple(
            dev_xvals->end(),
            dev_yvals->end())),
        functor,
        initVal,
        thrust::plus<double>());
}
```

GNU Compiler vs. Intel Compiler (cont.)




OpenMP Backend on Intel Xeon Phi (Native Execution)

Compile (Gauss fit: Native Xeon Phi)

- Gauss Fit
 - `icc solution2.cpp -O2 -o omp.out -mmic -openmp -DTHRUST_DEVICE_BACKEND=THRUST_DEVICE_BACKEND_OMP -liomp5 -I $CUDA_HOME/include -L~/local/mic/lib -L./rootstuff/ -lRootUtils -Wl,-rpath=/nfs/17/osc0724/local/mic/lib -Wl,-rpath=/nfs/17/osc0724/mic/nativeMic/solution2/rootstuff`
- Rootstuff
 - `icc -l. -O2 -mmic -m64 -fPIC -pthread -g -c -o TMinuit.o TMinuit.cc`
 - `icc -l. -O2 -mmic -m64 -fPIC -pthread -g -c -o TRandom.o TRandom.cc`
 - `icc -l. -O2 -mmic -m64 -fPIC -pthread -g -c -o TRandom3.o TRandom3.cc`
 - `icc -mmic -shared -Wl,-soname,libRootUtils.so -O2 -mmic -m64 -fPIC -pthread -g -o libRootUtils.so TMinuit.o TRandom.o TRandom3.o\`


Gauss fit: Native Xeon Phi (240 Threads)

Gauss fit Xeon Phi (240 Threads)




Gauss fit: Native Xeon Phi (180 Threads)

Gauss fit Xeon Phi (180 Threads)




Gauss fit: Native Xeon Phi (120 Threads)

Gauss fit Xeon Phi (120 Threads)

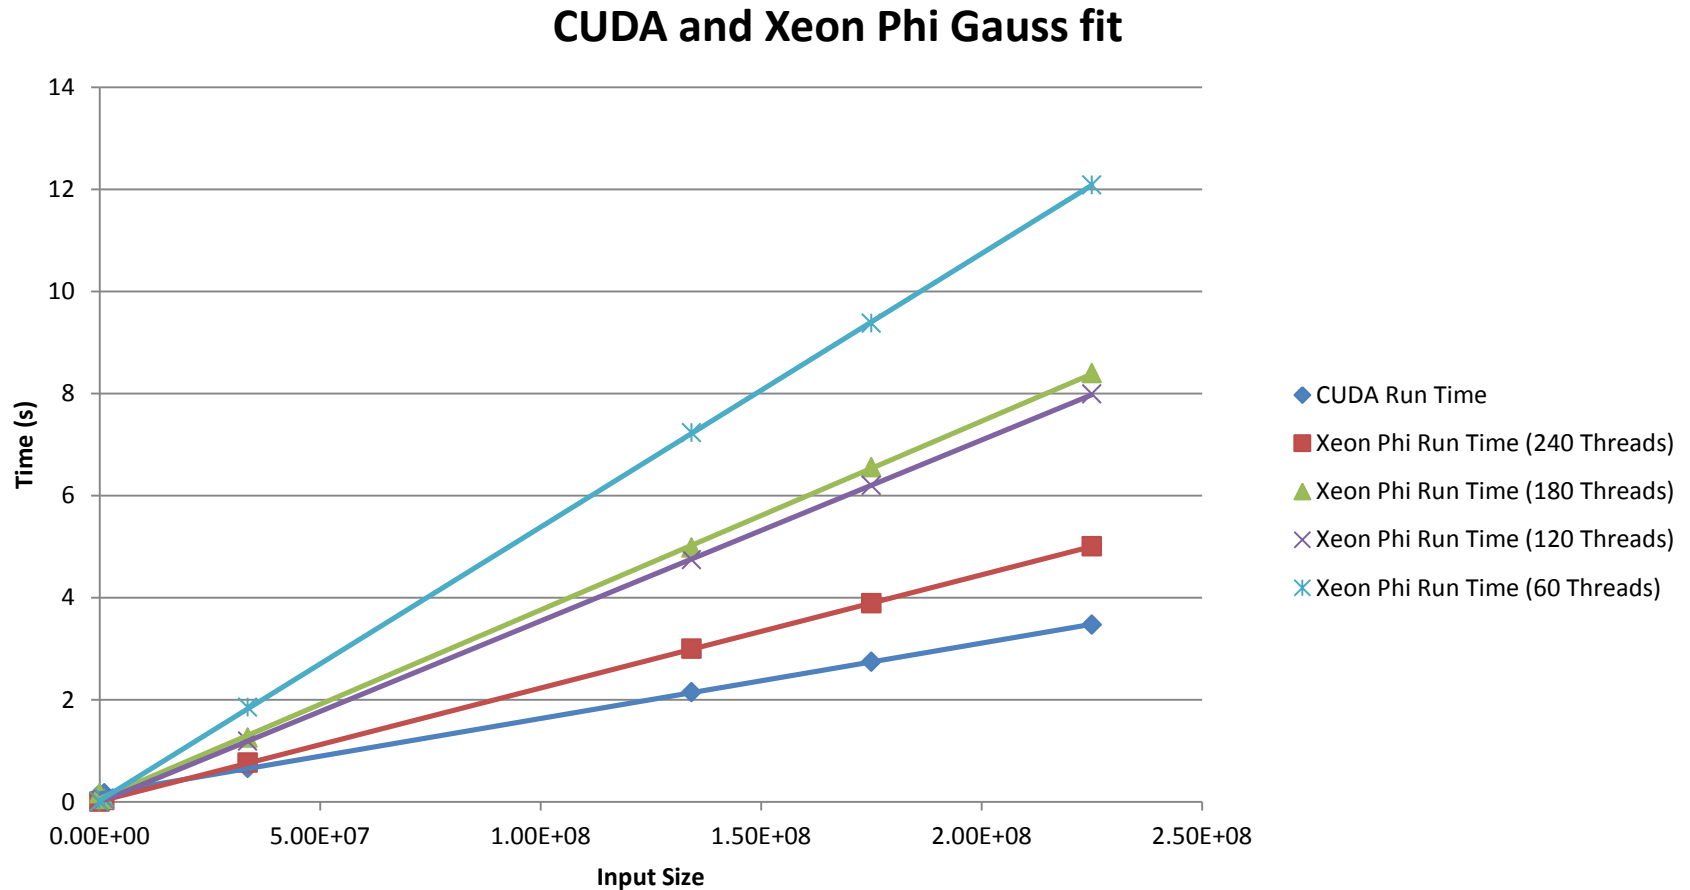


Gauss fit: Native Xeon Phi (60 Threads)

Gauss fit Xeon Phi (60 Threads)



Gauss fit: CUDA vs. Native Xeon Phi




Compile

(Chi-square fit: Native Xeon Phi)

- Chi-square Fit
 - `icc example2.cpp -O2 -o omp.out -mmic -openmp -DTHRUST_DEVICE_BACKEND=THRUST_DEVICE_BACKEND_OMP -liomp5 -I $CUDA_HOME/include -L~/local/mic/lib -L./rootstuff/ -lRootUtils -Wl,-rpath=/nfs/17/osc0724/local/mic/lib -Wl,-rpath=/nfs/17/osc0724/mic/nativeMic/example2/rootstuff`
- Rootstuff
 - `icc -l. -O2 -mmic -m64 -fPIC -pthread -g -c -o TMinuit.o TMinuit.cc`
 - `icc -l. -O2 -mmic -m64 -fPIC -pthread -g -c -o TRandom.o TRandom.cc`
 - `icc -l. -O2 -mmic -m64 -fPIC -pthread -g -c -o TRandom3.o TRandom3.cc`
 - `icc -mmic -shared -Wl,-soname,libRootUtils.so -O2 -mmic -m64 -fPIC -pthread -g -o libRootUtils.so TMinuit.o TRandom.o TRandom3.o\`


Chi-square fit: Native Xeon Phi (240 Threads)

Chi-square fit Xeon Phi (240 Threads)




Chi-square fit: Native Xeon Phi (180 Threads)

Chi-square fit Xeon Phi (180 Threads)




Chi-square fit: Native Xeon Phi (120 Threads)

Chi-square fit Xeon Phi (120 Threads)



Chi-square fit: Native Xeon Phi (60 Threads)

Chi-square fit Xeon Phi (60 Threads)



Chi-square fit: CUDA vs. Native Xeon Phi

CUDA and Xeon Phi Chi-square fit

